

# Big Data Processing, 2014/15

## **Lecture 6: MapReduce - behind the scenes continued (a very mixed bag)**

**Claudia Hauff (Web Information Systems)**  
**[ti2736b-ewi@tudelft.nl](mailto:ti2736b-ewi@tudelft.nl)**

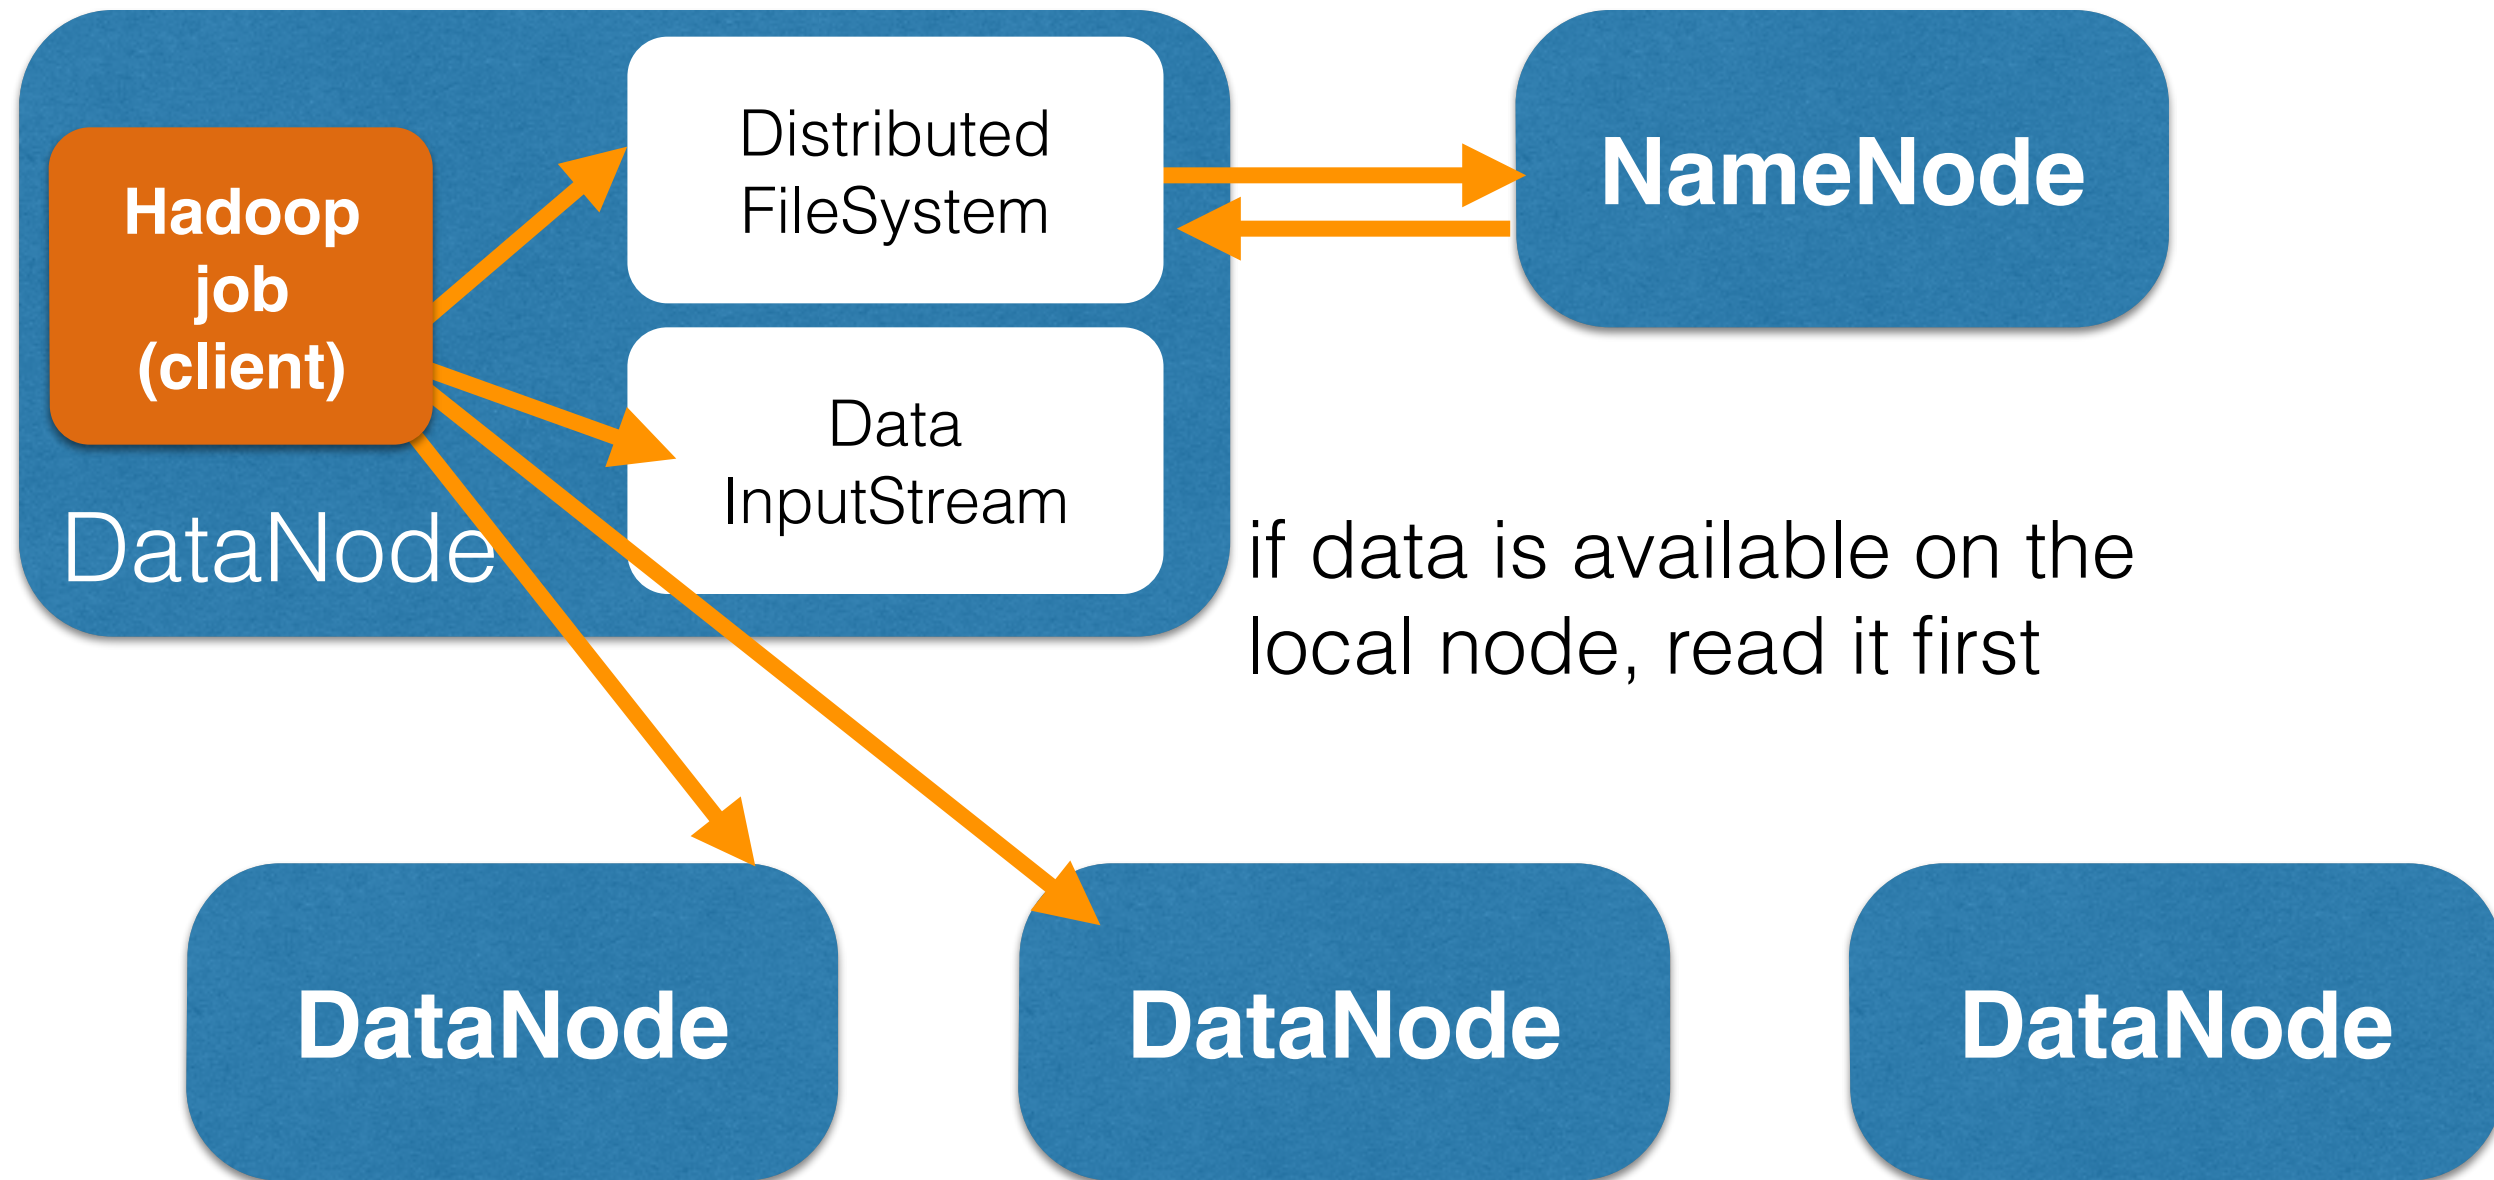
# Course content

- Introduction
- Data streams 1 & 2
- The MapReduce paradigm
- **Looking behind the scenes of MapReduce:** HDFS & **Scheduling**
- Algorithm design for MapReduce
- A high-level language for MapReduce: Pig 1 & 2
- MapReduce is not a database, but HBase nearly is
- Lets iterate a bit: Graph algorithms & Giraph
- How does all of this work together? ZooKeeper/Yarn


# Learning objectives

- **Exploit** Hadoop's Counters and setup/cleanup efficiently
- **Explain** how Hadoop addresses the problem of job scheduling
- **Explain** Hadoop's shuffle & sort phase and **use** that knowledge to improve your Hadoop code
- **Implement** strategies for efficient data input

# Question: what happens in each stage of the “read operation” ?



**Question: if the 3 map tasks started at the same time, in what order will they finish?**



Data center



**Question: which of the following GFS components reside on a chunkserver?**

**file**

**block**

**metadata**

**chunk**

**checksum**

**namespace**

**master**

**heartbeat**

**file permissions**

# Hadoop Programming

## Revisited: setup and cleanup

# Setup & cleanup

**Programmer  
“hints” the number  
of mappers to use**

- One `MAPPER` object for each map task
  - Associated with a sequence of key/value pairs (the “input split”)
  - `map( )` is called for each key/value pair by the execution framework
- One `REDUCER` object for each reduce task
  - `reduce( )` is called once per intermediate key
- `MAPPER/REDUCER` are Java objects -> allows side effects
  - Preserving state across multiple inputs
  - Initialise additional resources
  - Emit (intermediate) key/value pairs in one go

**Programmer can  
set the number of  
reducers**



# Setup

## Setup useful for one-off operations:

- opening an SQL connection
- loading a dictionary
- etc.

*WordCount\* - count only valid dictionary terms*

```
1 public class MyMapper extends
2     Mapper<Text, IntWritable, Text, IntWritable> {
3
4     private Set<String> dictionary; //all valid words
5
6     public void setup(Context context) throws IOException {
7         dictionary = Sets.newHashSet();
8         loadDictionary(); //define dictionary
9     }
10
11    public void map(Text key, IntWritable val, Context context)
12        throws IOException, InterruptedException {
13        if(!dictionary.contains(key.toString()))
14            return;
15        context.write(key, new IntWritable(1));
16    }
17 }
```

Called once in the life cycle of a Mapper object: before any calls to map()

Called once for each key/value pair that appears in the input split

# Cleanup

*WordCount\*\* - how many words start with the same letter?*

```
1 public class MyReducer extends
2   Reducer<PairOfIntString, FloatWritable, NullWritable, Text> {
3   private Map<Character, Integer> cache;
4
5   public void setup(Context context) throws IOException {
6       cache = Maps.newHashMap();
7   }
8   public void reduce(PairOfIntString key, Iterable<IntWritable>
9       values, Context context) throws
10      IOException, InterruptedException {
11       char c = key.toString().charAt(0);
12       for(IntWritable iw : values){
13           //add iw to the current value of key c in cache
14       }
15   }
17   public void cleanup(Context context) throws IOException,
18       InterruptedException {
19       for (Character c : cache.keySet()) {
20           context.write(new Text(c), new IntWritable(cache.get(c)));
21       }
22   }
23 }
```

# Cleanup

*WordCount\*\* - how many words start with the same letter?*

```

1 public class MyReducer extends
2   Reducer<PairOfIntString, FloatWritable, NullWritable, Text> {
3   private Map<Character, Integer> cache
4
5   public void setup(Context context) throws IOException, InterruptedException {
6       cache = Maps.newHashMap();
7   }
8   public void reduce(PairOfIntString key, Iterable<IntWritable>
9       values, Context context) throws IOException, InterruptedException {
10
11       char c = key.toString().charAt(0);
12       for(IntWritable iw : values){
13           //add iw to the current value
14       }
15   }
16
17   public void cleanup(Context context) throws IOException,
18       InterruptedException {
19       for (Character c : cache.keySet()) {
20           context.write(new Text(c), new FloatWritable(cache.get(c)));
21       }
22   }
23 }

```

Called once in the life cycle of a Reducer object: before any calls to `reduce()`

Called once for each key that was assigned to the reducer

Called once in the life cycle of a Reducer object: after all calls to `reduce()`

# Hadoop Programming

## Revisited: Counters

# Counter basics

- **Gathering data about the data** we are analysing, e.g.
  - Number of key/value pairs processed in map
  - Number of empty lines/invalid lines
- Wanted:
  - **Easy** to collect
  - **Viewable during job execution** (stop Hadoop job early at too many invalid key/value pairs)
- What about log messages?
  - Write to the error log when an invalid line occurs
  - Hadoop's logs are huge, you need to know where to look
  - Aggregating stats from the logs requires another pass over it



# Counter basics

- **Gathering data about the data** we are analysing, e.g.
  - Number of key/value pairs processed in map
  - Number of empty lines/invalid lines

**WordCount** example: what if we want to know more?

- How many words are not in the dictionary?
- How many words could not be parsed?
- How many words have less than two characters?

**Question: how can you achieve that with your current Hadoop knowledge?**



# Counter basics

- Counters: Hadoop's way of **aggregating** statistics
- Counters **count** (increment)
- **Built-in counters** maintain **metrics** of the job
  - MapReduce counters (e.g. #skipped records by all maps)
  - File system counters (e.g. #bytes read from HDFS)
  - Job counters (e.g. #launched map tasks)
- You have already seen them

# Question: what are the reasons for the discrepancy in the amount of data read and written?

- Counters: Hadoop's way of **aggregating** statistics

- Map-Reduce Framework

Map input records=5903

- Map output records=47102

Combine input records=47102

Combine output records=8380

Reduce output records=5934


File System Counters

FILE: Number of bytes read=118124

FILE: Number of bytes written=1075029

HDFS: Number of bytes read=996209

HDFS: Number of bytes written=59194



# Built-in vs. user-defined

- **Built-in counters:** maintained by the JobTracker
- **User-defined Counters** are maintained by the task with which they are associated
  - Periodically sent to the Tasktracker and then the Jobtracker for global aggregation

**Counter values are only definite once the job has completed (Counters may go down if a task fails!)**

# Code example

*WordCount\* - count words and chars*

```
1 enum Records {
2     WORDS, CHARS;
3 };
4 public class WordCount {
5     public static class Map extends MapReduceBase implements
6         Mapper<LongWritable, Text, Text, IntWritable> {
7
8         public void map(LongWritable key, Text value,
9             OutputCollector< Text, IntWritable> output,
10             Reporter reporter) throws IOException {
11             String[] tokens = value.toString().split(" ");
12
13             for (String s : tokens) {
14                 output.collect(new Text(s), new IntWritable(1));
15                 reporter.getCounter(Records.WORDS).increment(1);
16                 reporter.getCounter(Records.CHARS).increment(s.length());
17             }
18         }
19     }
```

several enum's possible: used to group counters

user-defined counters appear automatically in the final status output

# Code example

*WordCount\* - count words and chars*

```
1 enum Records
2     WORDS, C
3 };
4 public class
5     public sta
6
7
8     public v
9
10
11     Stri...
12
13     for
14         ou
15         re
16         re
17     }
18 }
19 }
```

## Map-Reduce Framework

Map input records=5903

Map output records=47102

Combine input records=47102

Combine output records=8380

Reduce output records=5934

## Records

CHARS=220986

WORDS=47102

atput,  
{

gth());

user-defined counters appear  
automatically in the final status output

# Code example II

```
1  enum Records { MAP_WORDS, REDUCE_WORDS; };
2
3  public class WordCount {
4      --> MAPPER
5      public void map(LongWritable key, Text value, OutputCollector<
6                      Text,IntWritable> output, Reporter reporter)
7                      throws IOException {
8
9          String[] tokens = value.toString().split(" ");
10         for (String s : tokens) {
11             output.collect(new Text(s), new IntWritable(1));
12             reporter.getCounter(Records.MAP_WORDS).increment(1);
13         }
14     }
15     --> REDUCER
16     public void reduce(Text key, Iterator<IntWritable> values,
17                       OutputCollector<Text,IntWritable> output,
18                       Reporter reporter) throws IOException {
19         int sum = 0;
20         while (values.hasNext())
21             sum += values.next().get();
22         reporter.getCounter(Records.REDUCE_WORDS).increment(sum);
23     }
24 }
```



# Question: Why does it make more sense in this scenario to define the Counter in the Mapper?

```
1  enum Records { MAP_WORDS, REDUCE_WORDS; };
2
3  public class WordCount {
4      --> MAPPER
5      public void map(LongWritable key, Text value, OutputCollector<
6                      Text,IntWritable> output, Reporter reporter)
7                      throws IOException {
8
9          String[] tokens = value.toString().split(" ");
10         for (String s : tokens) {
11             output.collect(new Text(s), new IntWritable(1));
12             reporter.getCounter(Records.MAP_WORDS).increment(1);
13         }
14     }
15     --> REDUCER
16     public void reduce(Text key, Iterator<IntWritable> values,
17                       OutputCollector<Text,IntWritable> output,
18                       Reporter reporter) throws IOException {
19         int sum = 0;
20         while (values.hasNext())
21             sum += values.next().get();
22         reporter.getCounter(Records.REDUCE_WORDS).increment(sum);
23     }
24 }
```

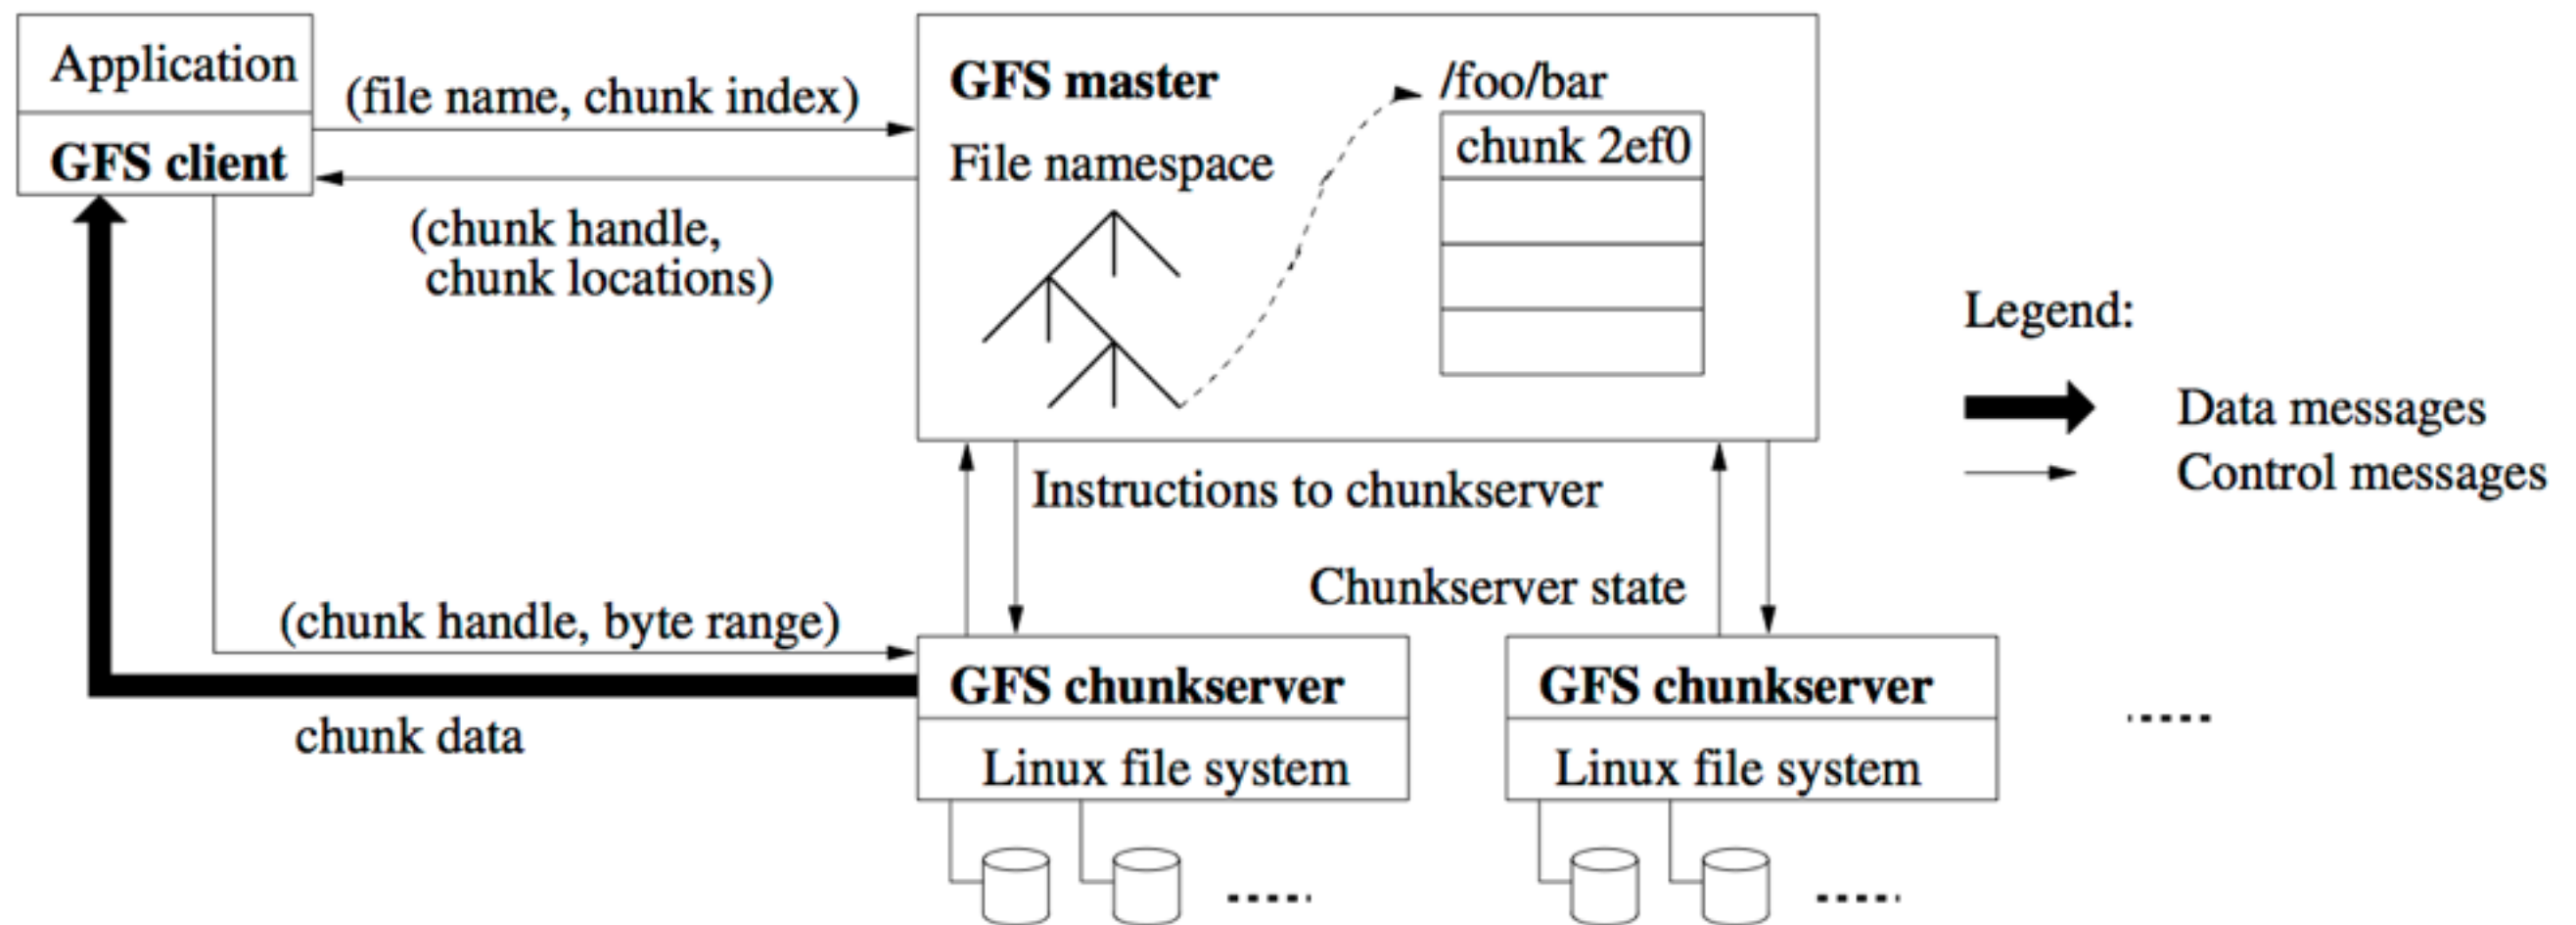
# Question: Does this code work as expected?

*WordCount across mapper/reducer*

```
1 class InMemoryCounter {
2     public int count;
3     public InMemoryCounter() {count=0;}
4 }
5 public class WordCount {
6     public static InMemoryCounter imc;
7     static { imc = new InMemoryCounter(); }
8     --> MAPPER
9     public void map(LongWritable key, Text value, OutputCollector<
10         Text, IntWritable> output, Reporter reporter) throws IOException {
11         String[] tokens = value.toString().split(" ");
12         for(String w : tokens) {
13             if(w.matches("[^a-zA-Z]")==true) //count non-alphanumeric terms
14                 imc.count++;
15             output.collect(new Text(w), new IntWritable(1));
16         }
17     --> REDUCER
18     public void reduce(Text key, Iterator<IntWritable> values,
19         OutputCollector<Text, IntWritable> output, Reporter reporter)
20         throws IOException {
21         while (values.hasNext()) {
22             int v = values.next().get();
23             if(key.toString().matches("[^a-zA-Z]")==false)//count the rest
24                 imc.count+=v;
25         }
26     }
27 }
```

# Job Scheduling

# Last time ... GFS/HDFS



**distributed file system:** file systems that manage the storage across a network of machines.

# What about the jobs?

- “Hadoop job”: unit of work to be performed (by a client)
  - Input data
  - MapReduce program
  - Configuration information
- Hadoop divides input data into **fixed size input splits**
  - One map task per split
  - One map function call for each record in the split
  - Splits are processed in parallel (if enough DataNodes exist)
- Job execution controlled by **JobTracker** and **TaskTrackers** (pre-YARN setup)

# What about the jobs?

**Question: What is the optimal input split size?**

- Configuration information

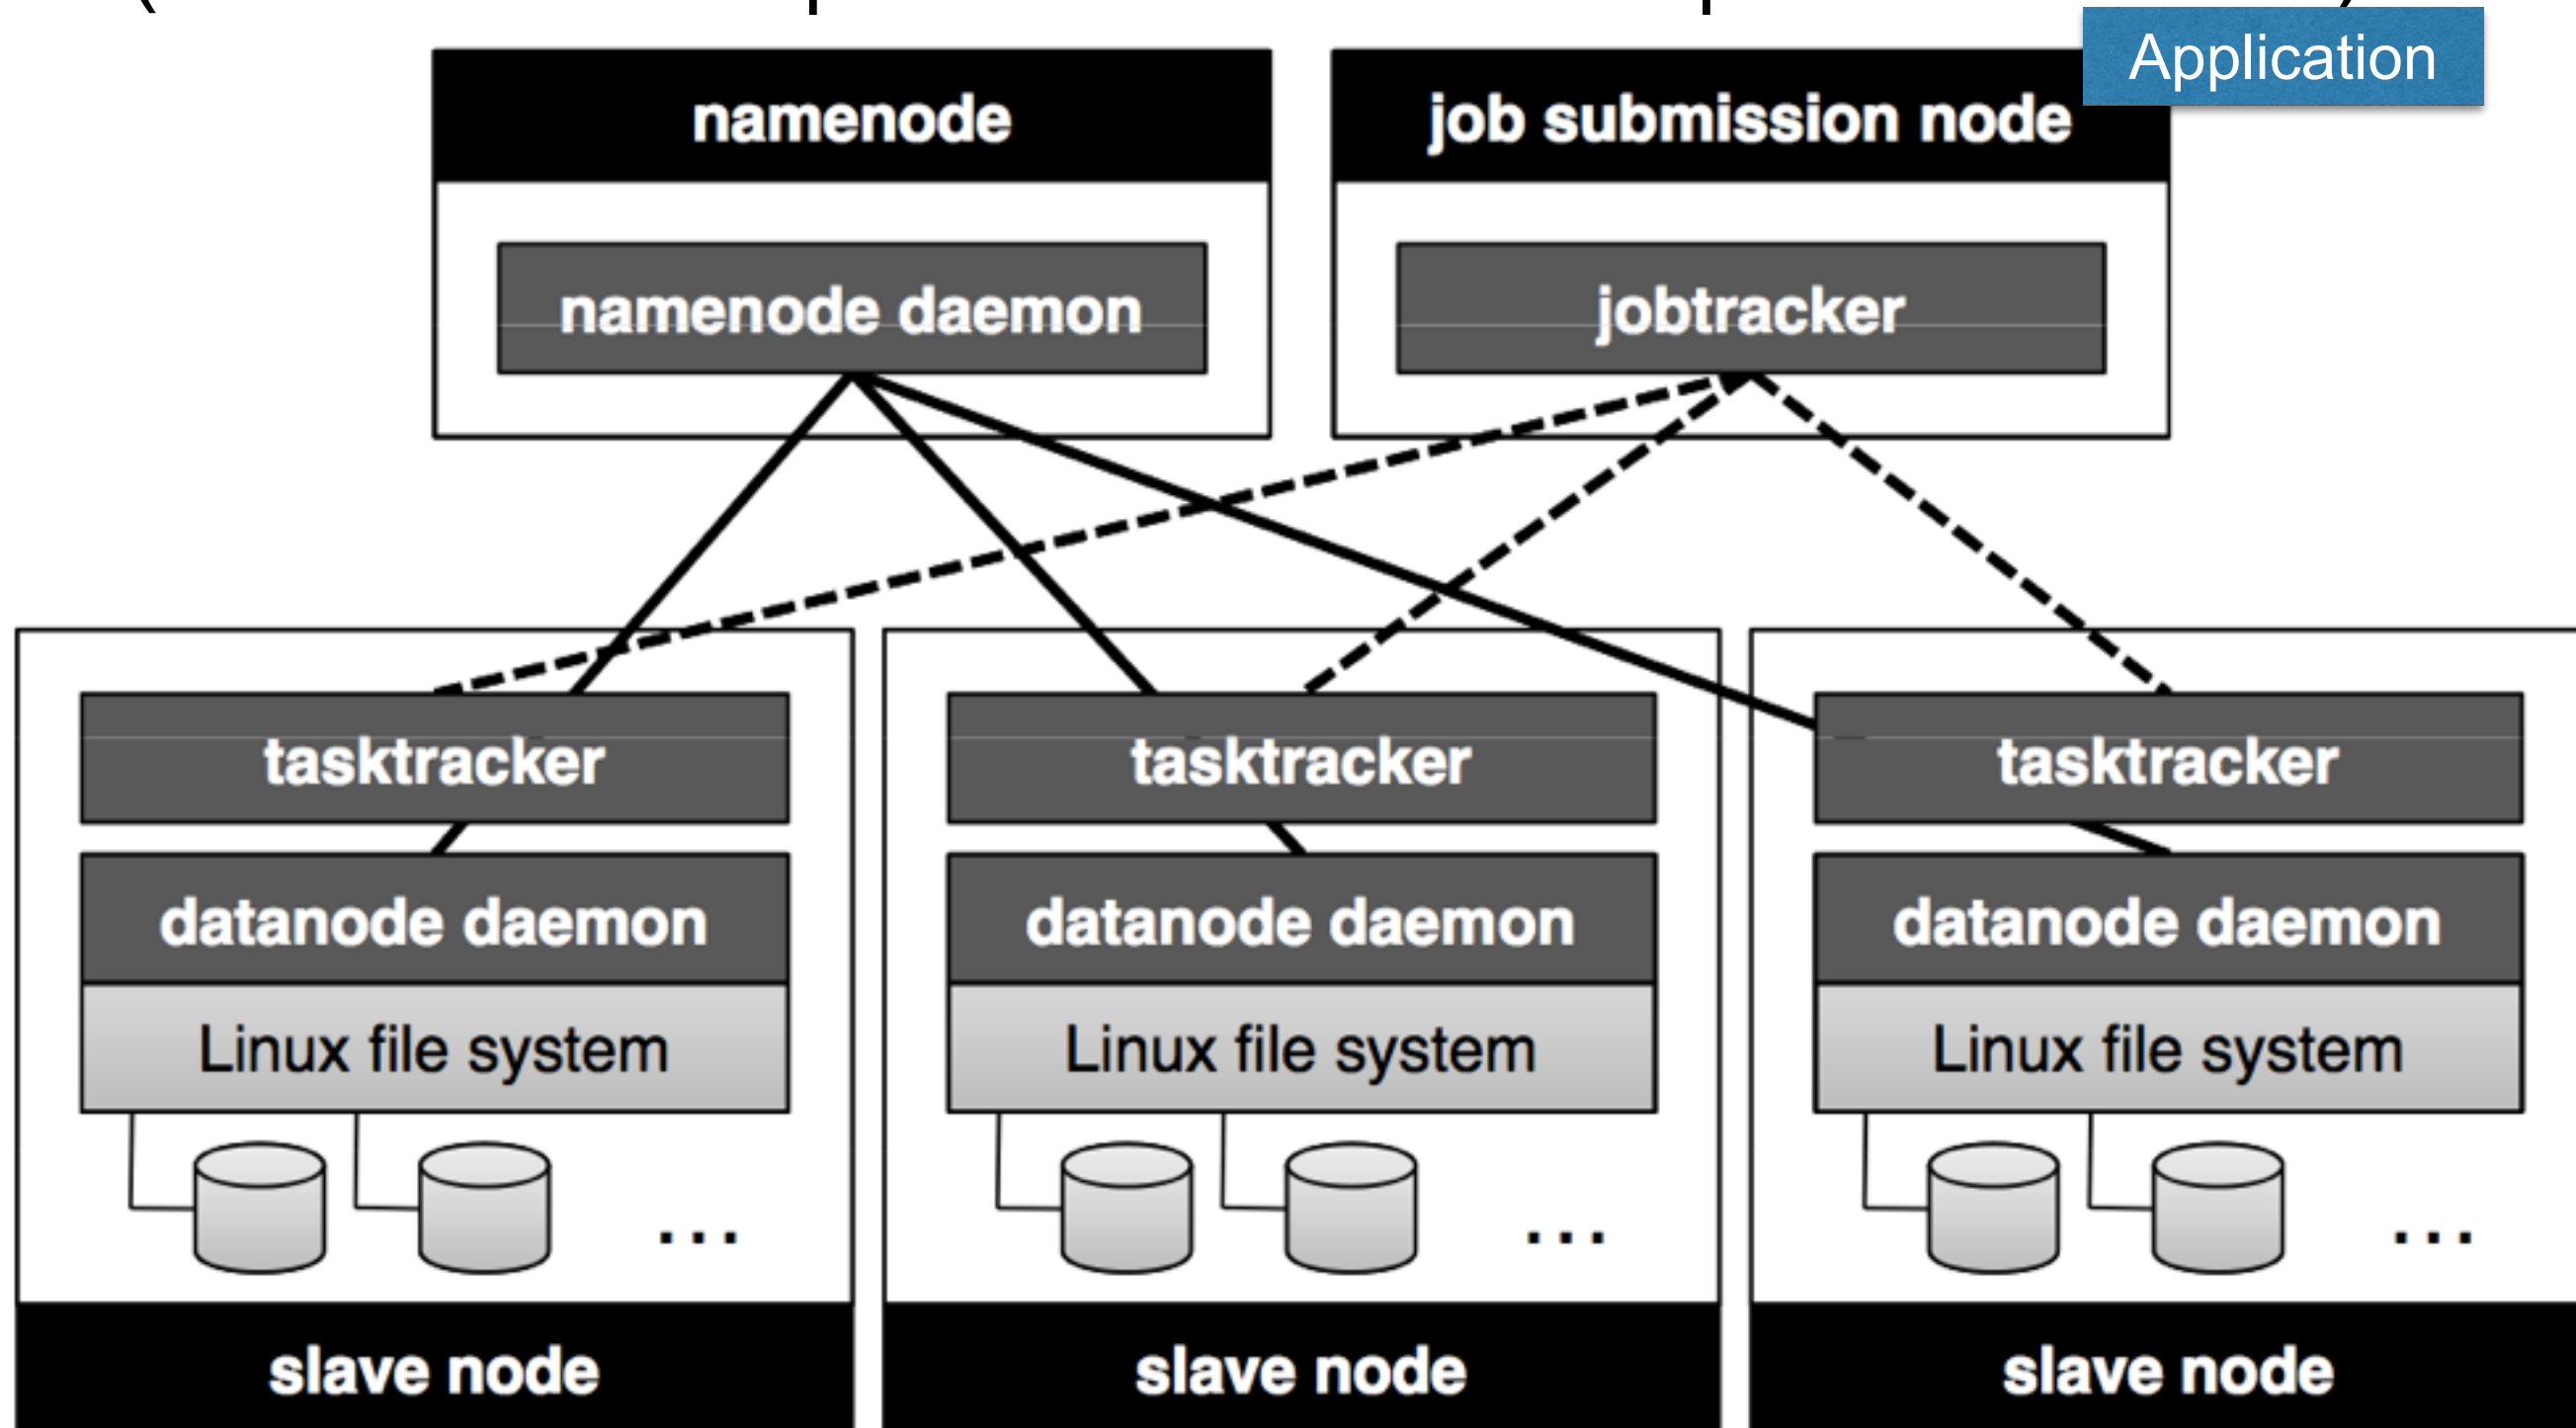
**Question: Can we also exploit data locality in the reducer?**

- Splits are processed in parallel (if enough DataNodes exist)
- Job execution controlled by **JobTracker** and **TaskTrackers** (pre-YARN setup)



# JobTracker and TaskTracker

## (Classic MapReduce or MapReduce 1)



# JobTracker and TaskTracker

## (Classic MapReduce or MapReduce 1)

- JobTracker
  - **One** JobTracker **per Hadoop cluster**
  - **Middleman** between your application and Hadoop (single point of contact)
  - Determines the **execution plan** for the application (files to process, assignment of nodes to tasks, task monitoring)
  - Takes care of (supposed) **task failures**
- TaskTracker
  - **One** TaskTracker **per DataNode**
  - Manages individual tasks
  - **Keeps in touch** with the JobTracker (via HeartBeats) - sends progress report & signals empty task slots

# YARN (MapReduce 2)

- JobTracker/TaskTrackers setup becomes a **bottleneck** in clusters with thousands of nodes
- As answer YARN has been developed (**Y**et **A**nother **R**esource **N**egotiator)
- YARN splits the JobTracker's tasks (job scheduling and task progress monitoring) into two daemons:
  - **Resource manager** (RM)
  - **Application master** (negotiates with RM for cluster resources; each Hadoop job has a dedicated master)

# Job scheduling

- Thousands of tasks may make up one job
- Number of tasks can exceed number of tasks that can run concurrently
  - Scheduler maintains task queue and tracks progress of running tasks
  - Waiting tasks are assigned nodes as they become available
- “Move code to data”
  - Scheduler starts tasks on node that holds a particular block of data needed by the task if possible

# Job scheduling

- Early on: **FIFO scheduler**
  - Job occupies the whole cluster while the rest waits
  - Not feasible in larger clusters
- Improvement: different job priorities VERY\_HIGH, HIGH, NORMAL, LOW, or VERY\_LOW
  - Next job is the one with the highest priority
  - No pre-emption: if a low priority job is occupying the cluster, the high priority job still has to wait
- Now: Fair Scheduler & Capacity Scheduler

# Fair Scheduler

- Goal: every user receives a **fair share** of the cluster capacity over time
- If a single job runs, it uses the entire cluster
  - As more jobs are submitted, free task slots are given away such that each user receives a “fair share”
  - Short jobs complete in reasonable time, long jobs keep progressing
- A user who submits more jobs than a second user will not get more cluster resources on average



# Fair Scheduler

- Jobs are placed in pools, default: one pool per user
- **Pre-emption**: if a pool has not received its fair share for a certain period of time, the scheduler will kill tasks in pools running over capacity to give more slots to the pool running under capacity
  - **Task kill != Job kill**
  - Scheduler needs to keep track of all users, resources used

# Capacity Scheduler

- Cluster is made up of a number of queues (similar to the Fair Scheduler pools)
- Each queue has an allocated capacity
- Within each queue, jobs are scheduled using FIFO with priorities
- Idea: users (defined using queues) simulate a separate MapReduce cluster with FIFO scheduling for each user

# Speculative execution

- Map phase is only as fast as slowest MAPPER
- Reduce phase is only as fast as slowest REDUCER
- Hadoop job is sensitive to stragglers (tasks that take unusually long to complete)
- Idea: identical copy of task executed on a second node; the output of whichever node finishes first is used (improvements up to 40%)
  - running task is killed
- Can be done for both MAPPER/REDUCER
- Strategy does not help if straggler due to skewed data distribution

# Speculative execution

**Question: Why is speculative execution in practice mostly restricted to map tasks?**

- Hadoop job is sensitive to stragglers (tasks that take unusually long to complete)

**Question: Can we use the Partitioner to avoid a skewed distribution (e.g. on WordCount)?**

- running task is killed
- Can be done for both MAPPER/REDUCER
- Strategy does not help if straggler due to skewed data distribution

# Shuffle & Sort

# Shuffle & sort phase

- **Hadoop guarantee:** the input to every reducer is sorted by key
- **Shuffle:** sorting of **intermediate key/value pairs** and transferring them to the reducers (as input)
- “Shuffle is the heart of MapReduce”
- Understanding shuffle & sort is vital to recognise job bottlenecks
- Disclaimer: constantly evolving (*again*), description most valid for Hadoop 0.2X



# A high-level view



## MAP TASK



input  
split

map()

in-memory  
buffer




partition,  
sort, and  
spill to disk

merge (disk)




reduce tasks



reduce()

## REDUCE TASK




sort phase



copy across  
the network

map tasks

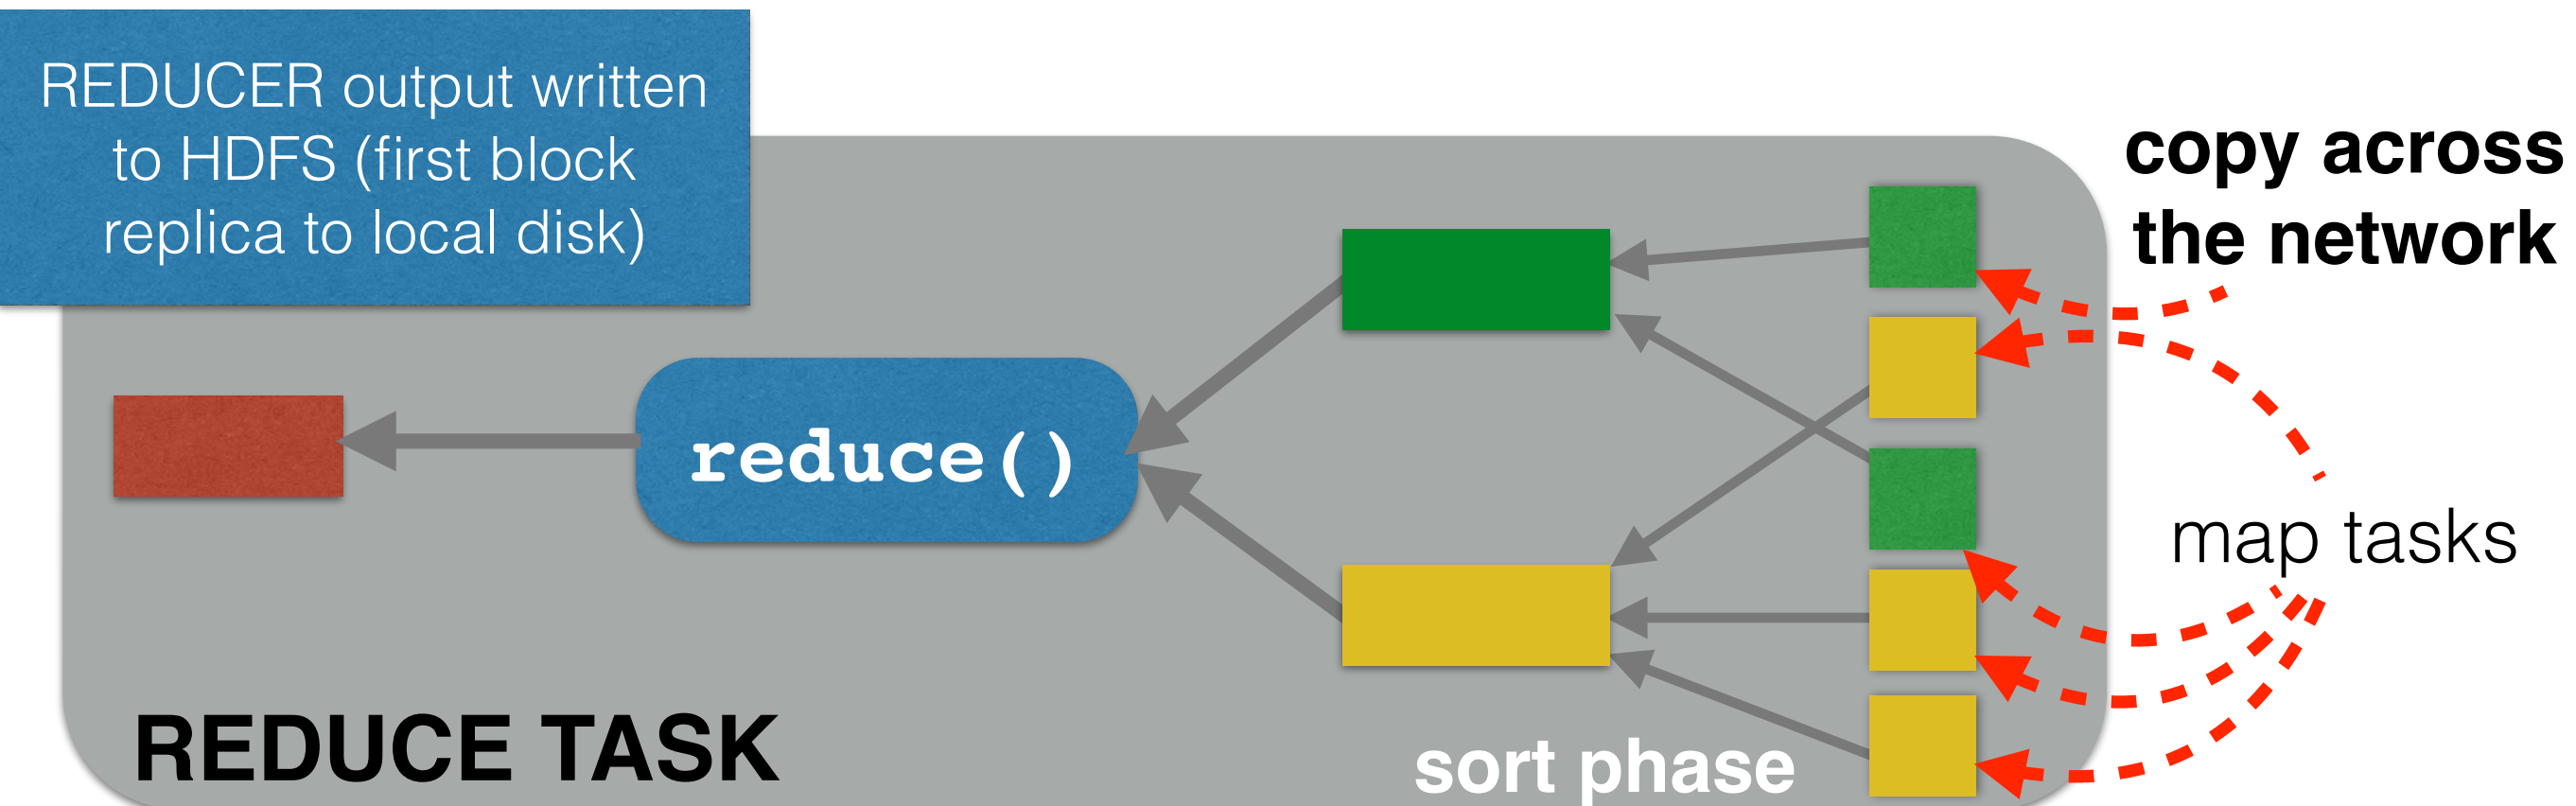
# Map side



- Map task writes output to memory buffer
- Once the buffer is full, a background thread spills the content to disk (spill file)
  - Data is partitioned corresponding to reducers they will be send to
  - Within partition, in-memory sort by key [combiner runs on the output of sort]
- After last `map()` call, the spill files are merged [combiner may run again]

# Reduce side

- Reducer requires the map output for its partition from **all map tasks of the cluster**
- Reducer starts copying data as soon as a map task completes ("copy phase")
- Direct copy to reducer's memory if the output is small, otherwise copy to disk
- In-memory buffer is merged and spilled to disk once it grows too large
- **Combiner may run again**
- Once **all** intermediate keys are copied the "sort phase" begins: merge of map outputs, maintaining their sort ordering



# A few more details

## MAP TASK



### What happens to the data written to local disk by the Mapper?

Jobtracker gives the signal for deletion after successful completion of the job.

### General rule for memory usage: map/reduce/shuffle

Shuffle should get as much memory as possible; write map/reduce with low memory usage (single spill would be best)

### How does the Reducer know where to get the data from?

- Successful map task informs task tracker which informs the job tracker (via heartbeat)
- Reducer periodically queries the job tracker for map output hosts until it has retrieved all of data

## REDUCE TASK



reduce tasks

copy across  
the network

# Sort phase recap

- Involves all nodes that executed map tasks and will execute reduce tasks
  - Job with  $m$  mappers and  $r$  reducers involves up to  $mr$  distinct copy operations
- Reducers can only start calling `reduce()` after all mappers are finished
  - **Key/value guarantee:** one key has all values “attached”
- Copying can start earlier for intermediate keys

# Summary

- Hadoop Counters, setup/cleanup
- Job scheduling
- Shuffle & sort



THE END