

# Big Data Processing, 2014/15

## Lecture 8: Pig Latin

**Claudia Hauff (Web Information Systems)**  
**[ti2736b-ewi@tudelft.nl](mailto:ti2736b-ewi@tudelft.nl)**

Start up CDH to follow along the Pig script examples. The example data used in the lecture is available on BB: **[TI2736-B]/Lectures/Lecture8/example\_data\_lecture8**

# Course content

- Introduction
- Data streams 1 & 2
- The MapReduce paradigm
- Looking behind the scenes of MapReduce: HDFS & Scheduling
- Algorithm design for MapReduce
- **A high-level language for MapReduce: Pig Latin 1 & 2**
- MapReduce is not a database, but HBase nearly is
- Lets iterate a bit: Graph algorithms & Giraph
- How does all of this work together? ZooKeeper/Yarn

# Learning objective

- **Translate** basic problems (suitable for MapReduce) into Pig Latin based on built-in operators

# Last time ...

Relation *Hyperlinks*

FROM	TO
url1	url2
url2	url3
url3	url5

- Database tables can be written out to file, one tuple per line
- MapReduce jobs can perform standard database operations
  - Most useful for operations that pass over most (all) tuples
- Joins are particularly tedious to implement in “plain” Hadoop

We don't just have  
joins .....

**No Pig Latin yet**

# Selections

Web\_pages

Url	Category	Last_crawl_ date	Page_length	Lng
news.yahoo.de	news	03-12-2013 07:08:45	765443	GER
nu.nl	news	03-12-2013 11:45:00	64435	NL
chess.com	game	23-10-2013 19:34:01	1264	EN
www.bbc.com/ sport/0/football/	sports	03-12-2013 14:13:22	6324	EN

**Question:** how can you do a selection in Hadoop?

# Projections

Web\_pages

Url	Category	Last_crawl_ date	Page_length	Lng
news.yahoo.de	news	03-12-2013 07:08:45	765443	GER
nu.nl	news	03-12-2013 11:45:00	64435	NL
chess.com	game	23-10-2013 19:34:01	1264	EN
www.bbc.com/ sport/0/football/	sports	03-12-2013 14:13:22	6324	EN

**Question:** how can you do a projection in Hadoop?



# Union

Web\_pages\_crawler1

Url	Category	Page_length	Lng
news.yahoo.de	news	765443	GER
nu.nl	news	64435	
chess.com	game	1264	
www.bbc.com/ sport/0/football/	sports	6324	

Web\_pages\_crawler2

Url	Category	Page_length	Lng
news.yahoo.de	news	765443	GER
volkskrant.nl	news	234445	NL
chessbase.com	game	1264	EN
www.bbc.com/ sport/0/football/	sports	6324	EN

**Question:** how can you do a union in Hadoop?



# Intersection

Web\_pages\_crawler1

Url	Category	Page_length	Lng
news.yahoo.de	news	765443	GER
nu.nl	news	64435	
chess.com	game	1264	
www.bbc.com/ sport/0/football/	sports	6324	

Web\_pages\_crawler2

Url	Category	Page_length	Lng
news.yahoo.de	news	765443	GER
volkskrant.nl	news	234445	NL
chessbase.com	game	1264	EN
www.bbc.com/ sport/0/football/	sports	6324	EN

**Question:** how can you do an intersection in Hadoop?

And now .... Pig Latin

# Pig vs. Pig Latin

- Pig: an **engine** for executing **data flows** in parallel on Hadoop.
- **Pig Latin**: the language for expressing data flows
- Pig Latin contains common data processing operators (**join**, **sort**, **filter**, ...)
- **User defined functions** (UDFs): developers can write their own functions to read/process/store the data


Pig is part of the CDH!!

# Pig on Hadoop

- Makes use of **HDFS** and the **MapReduce core** of Hadoop
  - By default, reads input from & writes output to HDFS
- Pig Latin scripts are **compiled** into **one or more** Hadoop jobs which are executed in order
- Pig Latin users need **not** to be aware of the algorithmic details in the map/shuffle/reduce phases
  - Pig **decomposes** operations into the appropriate map and/or map/reduce phases **automatically**

# Pig Latin

- A parallel **dataflow language**: users describe **how** data is read, processed and stored
  - Dataflows can be simple (e.g. “counting words”) or complex (multiple inputs are joined, data is split up into streams and processed separately)
  - Formally: a Pig Latin script describes a **Directed Acyclic Graph**
- directed graph, no directed cycles



# Pig vs. OO & SQL

- OO programming languages describe control flow with data flow as side effect, Pig Latin describes data flow (**no control constructs** such as `if`)

Pig	SQL
<b>Procedural</b> : script describes <b>how</b> to process the data	<b>Descriptive</b> : query describes <b>what</b> the output should be
Workflows can contain <b>many</b> data processing operations	<b>One</b> query answers one question (*subqueries)
Schemas may be <b>unknown or inconsistent</b>	RDBMSs have <b>defined</b> schemas
Reads files from HDFS (and other sources)	Data is read from database tables



# Pig vs. Hadoop

Pig	Hadoop
Standard data-processing operations are built-in ( <a href="#">filter</a> , <a href="#">join</a> , <a href="#">group-by</a> , <a href="#">order-by</a> , ...)	<a href="#">Group-by</a> and <a href="#">order-by</a> exist. <a href="#">Filtering</a> and <a href="#">projection</a> are easy to implement. <a href="#">Joins</a> are hard work
Contains non-trivial implementations of data operators (e.g. for skewed key distributions reducer load can be rebalanced)	Load re-balancing based on key/value distributions not available
<a href="#">Error checking and optimization</a>	Code within map & reduce is <a href="#">executed as-is</a>
Pig Latin scripts are <a href="#">easy to understand, maintain and extend</a>	Relatively <a href="#">opaque</a> code with a lot of (ever changing) boilerplate
<a href="#">Few lines of code and a short development time</a>	A large amount of boilerplate

# Pig vs. Hadoop

Pig	Hadoop
Standard data-processing operations are built-in ( <a href="#">filter</a> , <a href="#">join</a> , <a href="#">group-by</a> , <a href="#">order-by</a> , ...)	<a href="#">Group-by</a> and <a href="#">order-by</a> exist. <a href="#">Filtering</a> and <a href="#">projection</a> are easy to implement. <a href="#">Joins</a> are hard work
Contains non-trivial implementations of data operators (e.g. for skewed key distributions reducer load can be rebalanced)	Load re-balancing based on key/value distributions not available

## Why then use Hadoop at all?

Pig heavily optimises **standard data operations**. Less common operations can be difficult to implement as Pig Latin is more restrictive than Hadoop.

# PigMix: Pig script benchmarks

A set of queries to test Pig's performance: how well does a Pig script perform compared to a direct Hadoop implementation?

Run date: August 27, 2009, run against top of trunk as of that day.

Pig 0.12 (4/4/2013)

Test	Pig run time	Java run time	Multiplier
PigMix_1	218	133.33	1.635
PigMix_2	99.333	48	2.07
PigMix_3	272	127.67	2.13
PigMix_4	142.33	76.333	1.87
PigMix_5	127.33	107.33	1.19
PigMix_6	135.67	73	1.86
PigMix_7	124.67	78.333	1.59
PigMix_8	117.33	68	1.73

Test	Pig run time	Java run time	Multiplier
PigMix_1	168	142	1.1830985915493
PigMix_2	71	62	1.14516129032258
PigMix_3	141	158	0.892405063291139
PigMix_4	93	87	1.06896551724138
PigMix_5	87	158	0.550632911392405
PigMix_6	93	81	1.14814814814815
PigMix_7	77	87	0.885057471264368
PigMix_8	62	57	1.08771929824561

# PigMix: Pig script benchmarks

A set of queries to test Pig's performance: how well does a Pig script perform compared to a direct Hadoop implementation?

anti-join:

```
SELECT
*
FROM table1 t1
LEFT JOIN table2 t2 ON t1.id = t2.id
WHERE t2.id IS NULL
```

Run date: August 27, 2009, run again

Test	Pig run time	Java run time	Multiplier
PigMix_1	218	133.333	30985915493
PigMix_2	99.333	48	516129032258
PigMix_3	272	127.67	2.13
PigMix_4	142.33	76.333	1.87
PigMix_5	127.33	107.33	1.19
PigMix_6	135.67	73	1.86
PigMix_7	124.67	78.333	1.59
PigMix_8	117.33	68	1.73

# Pig is useful for

- **ETL** (extract transform load) data pipelines
  - Example: web server logs that need to be cleaned before being stored in a data warehouse
- Research on raw data
  - Pig **handles erroneous**/corrupt **data** entries gracefully (cleaning step can be skipped)
  - Schema can be **inconsistent** or missing
  - Exploratory analysis can be performed **quickly**
- Batch processing
  - Pig Latin scripts are internally **converted to Hadoop jobs** (the same advantages/disadvantages apply)

# Pig philosophy

- **Pigs eat anything**
  - Pig operates on any data (schema or not, files or not, nested or not)
- **Pigs live anywhere**
  - Parallel data processing language; implemented on Hadoop but not tied to it
- **Pigs are domestic animals**
  - Easily controlled and modified
- **Pigs fly**
  - Fast processing



# History of Pig

- Research project at **Yahoo! Research**
- Paper about Pig prototype published in **2008**
- **Motivation:**
  - Data scientists spent too much time writing Hadoop jobs and not enough time analysing the data
  - Most Hadoop users know SQL well
- **Apache top-level project** in 2010

# Pig's version of WordCount

A screencast explaining the code line by line is available on Blackboard!

**TI2736-B/Lectures/Lecture8/Screencast: first Pig example**

```
-- read the file pg46.txt line by line, call each record line
cur = load 'pg46.txt' as (line);

-- tokenize each line, each term is now a record called word
words = foreach cur generate flatten(TOKENIZE(line)) as word;

-- group all words together by word
grp = group words by word;

-- count the words
cntd = foreach grp generate group, COUNT(words);

/*
 * start the Hadoop job and print results
 */
dump cntd;
```

**5 lines of code in Pig vs. 50 in plain Hadoop**

# Pig's version of V

```
-- read the file pg46.txt line by line,
christmas_book = load 'pg46.txt' as (li

-- tokenize each line, each term is now
words = foreach input generate flatten(

-- group all words together by word
grp = group words by word;

-- count the words
cntd = foreach grp generate group, CO

/*
 * start the Hadoop job and print result
 */
dump cntd;
```

5 lines of code in

```
(well-remembered,1)
(wine-merchant's,2)
(blindman's-buff.,1)
(entered--flushed,1)
(extinguisher-cap,1)
(highly-decorated,1)
(http://pglaf.org,2)
(including--which,1)
(knocker!--Here's,1)
(notwithstanding.,2)
(self-accusatory.,1)
(stagnant-blooded,1)
(unenforceability,1)
(fellow-'prentice.,1)
(fellow-passengers,1)
(five-and-sixpence,1)
(gbnewby@pglaf.org,1)
(sticking-plaister,1)
(strait-waistcoat.,1)
(surprised-looking,1)
(thread-the-needle,1)
(www.gutenberg.net,3)
(pleasantest-spoken,1)
(shabby--compounded,1)
(business@pglaf.org.,1)
(trademark/copyright,1)
(counting-house--mark,1)
(http://www.pglaf.org.,1)
(weathercock-surmounted,1)
(http://pglaf.org/donate,1)
(http://www.gutenberg.net,1)
(snowball--better-natured,1)
(http://gutenberg.net/license,1)
(http://pglaf.org/fundraising.,1)
(http://www.gutenberg.net/4/46/,1)
(,0)
grunt>
```

# Another example:

*Top clicked URL by users age 18-25*

John	18
Tom	24
Alfie	45
Ralf	56
Sara	19
Marge	27

**users:** name & age

John	url1
John	url2
Tom	url1
John	url2
Ralf	url4
Sara	url3
Sara	url2
Marge	url1

**clicks:** name & url

```
set io.sort.mb 5;
-- the top URL clicked by users age 18-25
users = load 'users' as (name,age);
filtered = filter users by age>=18 and age<=25;
clicks = load 'clicks' as (user,url);
joined = join filtered by name, clicks by user;

grouped = group joined by url;
summarized = foreach grouped generate group, COUNT(joined) as
    amount_clicked;
sorted = order summarized by amount_clicked desc;

top1 = limit sorted 1;
Store top1 into 'top1site';
```

A screencast explaining the code line by line is available on Blackboard!

**TI2736-B/Lectures/Lecture8/Screencast: top clicked URL**

# Pig is customisable

- All parts of the processing path are customizable
  - Loading
  - Storing
  - Filtering
  - Grouping
  - Joining
- Can be altered by **user-defined functions** (UDFs)

# Grunt: running Pig

- Pig's interactive shell

testing: local file system

real analysis: HDFS

- Grunt can be started in **local** and **MapReduce mode**

```
pig -x local      pig
```

Errors do not kill the chain of commands

- Useful for sampling data (a pig feature)
- Useful for prototyping: scripts can be entered interactively
  - Basic syntax and semantic checks (errors do not kill the chain of commands)
  - Pig executes the commands (starts a chain of Hadoop jobs) once **dump** or **store** are encountered



# Grunt: running Pig

- Pig's interactive shell

testing: local file system

real analysis: HDFS

- Grunt can be started in **local** and **MapReduce mode**

```
pig -x local      pig
```

Errors do not kill the chain of commands

- Useful for sampling data (a pig feature)
- Useful for prototyping: scripts can be entered interactively
  - Basic syntax and semantic checks (errors do not kill the chain of commands)
  - Pig executes the command (jobs) once **dump** or **store**

Other ways of running Pig Latin:

(1) `pig script.pig`

(2) Embedded in Java programs  
(`PigServer` class)

null: value  
unknown  
(SQL-like)

# Pig's data model

`java.lang.String`

- Scalar types: `int`, `long`, `float`, `double`, `chararray`, `bytearray`

`DataByteArray`, wraps `byte[]`

- Three complex types that can contain data of any type (nested)
  - **Maps**: `chararray` to data element mapping (values can be of different types) `[name#John,phone#5551212]`
  - **Tuples**: ordered collection of Pig data elements; tuples are divided into fields; analogous to rows (tuples) and columns (fields) in database tables `(John,18,4.0F)`
  - **Bags**: unordered collection of tuples (tuples cannot be referenced by position) `{(bob,21),(tim,19),(marge,21)}`

# Schemas

- Remember: pigs eat anything
- Runtime declaration of schemas
- Available schemas used for error-checking and optimization

```
[cloudera@localhost ~]$ pig -x local  
grunt> records = load 'table1' as (name:chararray, syear:chararray,  
>>grade:float);
```

Pig reads three fields per line, **truncates** the rest; **adds null** values for missing fields

**as** indicates the schema.

```
grunt> describe records;  
records: {name: chararray,syear: chararray,grade: float}
```

# Schemas

- What about data with 100s of columns of known type?
  - Painful to add by hand every time
  - Solution: store schema in metadata repository  
Apache HCatalog – Pig can communicate with it

table and storage management layer - offers a relational view of data in HDFS.

- Schemas are not necessary (but useful)

# A guessing game

column names, no types

```
[cloudera@localhost ~]$ pig -x local
grunt> records = load 'table1' as (name,syear,grade);
grunt> describe records;
records: {name: bytearray,syear: bytearray,grade: bytearray}
```

- Pig makes intelligent type guesses based on data usage (remember: nothing happens before we use the **dump/store** commands)
- If it is not possible to make a good guess, Pig uses the bytearray type (default type)

# Default names

column types, no names

```
grunt> records2 = load 'table1' as(chararray,chararray,float);  
grunt> describe records2;  
records2: {val_0: chararray, val_1: chararray, val_2: float}
```

- Pig assigns default names if none are provided
- Saves typing effort, makes complex programs difficult to understand ...



# No need to work with unwanted content

Read only the first column

```
grunt> records3 = load 'table1' as(name);  
grunt> dump records3;  
(bob)  
(jim)  
. . .
```

- We can select which file content we want to process

# More columns than data

```
grunt> records4 = load 'table1' as(name,syear,grade,city,bsn);
grunt> dump records4;
(bob,1st_year,8.5,,)
(jim,2nd_year,7.0,,)
(tom,3rd_year,5.5,,)
..
```

The file contains 3 “columns”  
– the remaining two columns  
are set to null

- Pig **does not throw an error** if the schema does not match the file content!
- Necessary for large-scale data where corrupted/incompatible entries are common
- Not so great for debugging purposes

# Pig: loading & storing

```
[cloudera@localhost ~]$ pig -x local
grunt> records = load 'table1' as (name:chararray,
>> syear:chararray, grade:float);
grunt> describe records;
records: {name: chararray,syear: chararray,grade: float}
grunt> dump records;
(bob,1st_year,8.5)
(jim,2nd_year,7.0)
(tom,3rd_year,5.5)
...
```

**relation consisting of tuples**

```
grunt> store records into 'stored_records' using PigStorage(',');
grunt> store records into 'stored_records2';
```

# Pig: loading & storing

tab separated text file

```
[cloudera@localhost ~]$ pig -x local
grunt> records = load 'table1' as (name:chararray,
>> syear:chararray, grade:float);
grunt> describe records;
records: {name: chararray, syear: chararray, grade: float}
grunt> dump records;
```

local file (URI)

```
(bob,1st_year,8.5)
(jim,2nd_year,7.0)
(tom,3rd_year,5.5)
...
```

dump runs a Hadoop job  
and writes output to screen

delimiter

**relation consisting of tuples**

```
grunt> store records into 'stored_records' using PigStorage(',');
grunt> store records into 'stored_records2';
```

default output is  
tab delimited

store runs a Hadoop  
job and writes output to

# Pig: loading and storing

```
[cloudera@localhost ~]$ ls stored_records/  
part-m-00000 SUCCESS  
[cloudera@localhost ~]$ more stored_records/part-m-00000  
bob,1st_year,8.5  
jim,2nd_year,7.0  
tom,3rd_year,5.5  
andy,2nd_year,6.0  
bob2,1st_year,7.5  
tim,2nd_year,8.0  
cindy,1st_year,8.5  
arie,2nd_year,6.5  
jane,1st_year,9.5  
tijs,1st_year,8.0  
claudia,2nd_year,7.5  
mary,3rd_year,9.5  
mark,3rd_year,8.5  
john,,  
ralf,,  
[cloudera@localhost ~]$ █
```

**store is a Hadoop job with  
only a map phase: part-m-\*\*\*\*\*  
(reducers output part-r-\*\*\*\*\*)**

# Relational operations

Transform the data by sorting, grouping, joining, projecting, and filtering.

# foreach

- Applies a set of expressions to every record in the pipeline
- Generates new records
- Equivalent to the projection operation in SQL

```
grunt> records = load 'table2' as (name,year,grade_1,grade_2);  
grunt> gradeless_records = foreach records generate name,year;  
grunt> gradeless_records = foreach records generate ..year;
```

```
grunt> diff_records = foreach records generate $3-$2,name;
```



# foreach

- Applies a set of expressions to every record in the pipeline

- Generates new records

- Equivalent to the projection operation in SQL

range of fields (useful when #fields is large)

```
grunt> records = foreach records generate name, year;  
grunt> records = foreach records generate ..year;  
grunt> 
```

```
grunt> diff_records = foreach records generate $3-$2, name;
```

fields can be accessed by their position

# foreach

- Evaluation function UDFs: take as input one record at a time and produce one output; Generates new records

```
grunt> records = load 'table1' as  
          (name:chararray,year:chararray,grade:float);  
grunt> grpd= group records by year;  
grunt> avgs = foreach grpd generate group, AVG(records.grade);  
grunt> dump avgs;  
(1st_year,8.4)  
(2nd_year,7.0)  
(3rd_year,7.83333333)  
(,)
```

Average: a built-in UDF

# filter

- Select records to keep in the data pipeline

```
grunt> filtered_records = FILTER records BY grade>6.5;
grunt> dump filtered_records;
(bob,1st_year,8.5)
(jim,2nd_year,7.0)
```

...

```
grunt> filtered_records = FILTER records BY grade>8 AND
      (year=='1st_year' OR year=='2nd_year');
```

```
grunt> dump filtered_records;
(bob,1st_year,8.5)
(cindy,1st_year,8.5)
```

...

```
grunt> notbob_records = FILTER records
      BY NOT name matches 'bob.*';
```

conditions can be combined

negation

regular expression

# filter

## inferred vs. defined data types

```
grunt> records = load 'table1' as (name,year,grade);
grunt> filtered_records = FILTER records BY grade>8
      AND (year=='1st_year' OR year=='2nd_year');
grunt> dump filtered_records;
```

inferred

```
grunt> records = load 'table1' as (name,year,grade);
grunt> filtered_records = FILTER records BY grade>8.0
      AND (year=='1st_year' OR year=='2nd_year');
grunt> dump filtered_records;
```

inferred

```
grunt> records = load 'table1' as
      (name:chararray,year:chararray,grade:f
grunt> filtered_records = FILTER records BY grade>8
      AND (year=='1st_year' OR year=='2nd_year');
grunt> dump filtered_records;
```

defined

A screencast explaining the code line by line is available on Blackboard!

**TI2736-B/Lectures/Lecture8/Screencast: inferred vs. defined**

# group

- Collect records together that have the **same key**

```
grunt> grouped_records = GROUP filtered_records BY syear;  
grunt> dump grouped_records;
```

```
(1st_year, {(bob, 1st_year, 8.5), (bob2, 1st_year, 7.5), (cindy,  
1st_year, 8.5), (jane, 1st_year, 9.5), (tijs, 1st_year, 8.0)})  
(2nd_year, {(tim, 2nd_year, 8.0), (claudia, 2nd_year, 7.5)})
```

two tuples, grouped together by the first field

bag of tuples,  
indicated by {}

```
grunt> describe grouped_records;
```

```
grunt> grouped_records: { filtered_records:  
{(name: chararray, syear: chararray, grade: float)}}}
```

name of grouping field




# group

Question: if the pipeline is in the map phase, what has to happen?

```
grunt> grouped_records = GROUP filtered_records BY syear;  
grunt> dump grouped_records;
```

Question: if the pipeline is in the reduce phase, what has to happen?

 two tuples, grouped together by the first field

 bag of tuples,  
indicated by {}

```
grunt> describe grouped_records;  
grunt> grouped_records: {  filtered_records:  
{(name: chararray, syear: chararray, grade: float)} }
```

 name of grouping field

# group

- There is no restriction on how many keys to group by
- All records with null keys end up in the same group

```
grunt> grouped_twice = GROUP records BY (year,grade);  
grunt> dump grouped_twice;
```

- In the underlying Hadoop job effects depend on phase:
  - Map phase: a reduce phase is enforced
  - Reduce phase: a map/shuffle/reduce is enforced



# group

- There is no restriction on how many keys to group by

- All records

```
grunt> gro  
grunt> dum
```

- In the un  
phase:

- Map p

```
((1st_year,7.5),{(bob2,1st_year,7.5)})  
((1st_year,8.0),{(tjjs,1st_year,8.0)})  
((1st_year,8.5),{(cindy,1st_year,8.5),(bob,1st_year,8.5)})  
((1st_year,9.5),{(jane,1st_year,9.5)})  
((2nd_year,6.0),{(andy,2nd_year,6.0)})  
((2nd_year,6.5),{(arie,2nd_year,6.5)})  
((2nd_year,7.0),{(jim,2nd_year,7.0)})  
((2nd_year,7.5),{(claudia,2nd_year,7.5)})  
((2nd_year,8.0),{(tim,2nd_year,8.0)})  
((3rd_year,5.5),{(tom,3rd_year,5.5)})  
((3rd_year,8.5),{(mark,3rd_year,8.5)})  
((3rd_year,9.5),{(mary,3rd_year,9.5)})  
((,),{(john,,),(ralf,,)})  
grunt> grouped_twice = group records by (year,grade);■
```

- Reduce phase: a map/shuffle/reduce is enforced

# order by

- Total ordering of the output data (including across partitions)
- Sorting according to the natural order of data types
- Sorting by maps, tuples or bags is not possible

```
grunt> records = load 'table1' as (name,year,grade);
grunt> graded = ORDER records BY grade,year;
grunt> dump graded;
(ralf,,)
(john,,)
. .
(tijs,1st_year,8.0)
(tim,2nd_year,8.0)
. .
```

The results are first ordered by grade and within tuples of the same grade also by year. Null values are ranked first.

# order by

- Pig balances the output across reducers
  1. Samples from the input of the order statement
  2. Based on the sample of the key distribution a “fair” partitioner is built

An additional Hadoop job for the sampling procedure is required.

Same key to different reducers!

- Example of sampled keys (3 reducers available):

a a a a c d x y z

{**a**, **(a,c,d)**, **(x,y,z)**}

# distinct

- Removes duplicate **records**

```
grunt> year_only = foreach records generate year;  
grunt> uniq_years = distinct year_only;  
(1st_year)  
(2nd_year)  
(3rd_year)  
( )
```

Works on entire records only,  
thus first a projection (line 1) is  
necessary.

Question: do we need a map and/or reduce phase here?

# join

- THE workhorse of data processing

```
grunt> records1 = load 'table1' as (name,year,grade);
grunt> records2 = load 'table3' as (name,year,country,km);
grunt> join_up = join records1 by (name,year),
                    records2 by (name,year);

grunt> dump join_up;
(jim,2nd_year,7.0,jim,2nd_year,Canada,164)
(tim,2nd_year,8.0,tim,2nd_year,Netherlands,)
. . .
```

- Pig also supports outer joins (values that do not have a match on the other side are included): left/right/full

```
grunt> join_up = join records1 by (name,year) left outer,
                    records2 by (name,year);
```



# join

- THE workhorses

```
grunt> records1 =  
grunt> records2 =  
grunt> join_up =
```

```
grunt> dump join_  
(jim,2nd_year,7.0  
(tim,2nd_year,8.0  
. . .
```

```
(bob,1st_year,8.5,,,) )  
(jim,2nd_year,7.0,jim,2nd_year,Canada,164)  
(tim,2nd_year,8.0,tim,2nd_year,Netherlands,)  
(tom,3rd_year,5.5,tom,3rd_year,Australia,6454)  
(andy,2nd_year,6.0,andy,2nd_year,Germany,445)  
(arie,2nd_year,6.5,,,) )  
(bob2,1st_year,7.5,bob2,1st_year,Belgium,12)  
(jane,1st_year,9.5,,,) )  
(john,,,,,,,,)  
(mark,3rd_year,8.5,,,) )  
(mary,3rd_year,9.5,,,) )  
(ralf,,,,,,,,)  
(tijs,1st_year,8.0,,,) )  
(cindy,1st_year,8.5,cindy,1st_year,Denmark,)  
(claudia,2nd_year,7.5,,,) )  
grunt> █
```

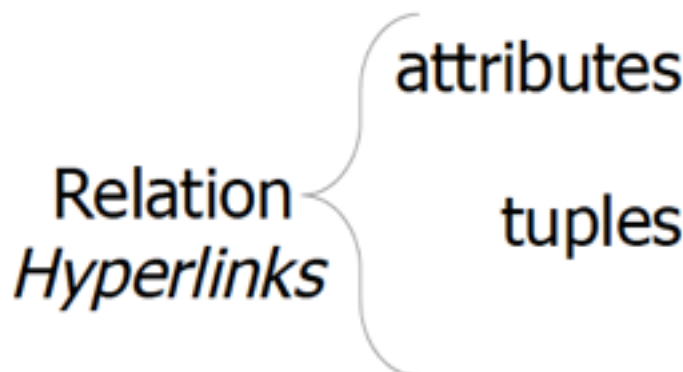
- Pig also supports outer joins (values that do not have a match on the other side are included): left/right/full

```
grunt> join_up = join records1 by (name,year) left outer,  
                      records2 by (name,year);
```

# join

- Self-joins are supported, though data needs to be loaded twice - very useful for graph processing problems

```
grunt> urls1 = load 'urls' as (A,B);
grunt> urls2 = load 'urls' as (C,D);
grunt> path_2 = join urls1 by B, urls2 by C;
grunt> dump path_2;
(url2,url1,url1,url2)
(url2,url1,url1,url4)
(url2,url1,url1,url3)
. . .
```



attributes	
FROM	TO
url1	url2
url2	url3
url3	url5

- Pig **assumes** that the **left** part of the join is the **smaller** data set



# limit

- Returns a limited number of records
- Requires a reduce phase to count together the number of records that need to be returned

```
grunt> urls1 = load 'urls' as (A,B);  
grunt> urls2 = load 'urls' as (C,D);  
grunt> path_2 = join urls1 by B, urls2 by C;  
grunt> first = limit path_2 1;  
grunt> dump first;  
(url2,url1,url1,url2)
```

- No ordering guarantees: every time limit is called it may return a different ordering

A sample command exists, e.g.  
`some = sample path 0.5;`  
to sample 50% of the data.

# illustrate

- Creating a sample data set from the complete one
  - Concise: small enough to be understandable to the developer
  - Complete: rich enough to cover all (or at least most) cases
- Random sample can be problematic for filter & join operations
- Output is easy to follow, allows programmers to gain insights into what the query is doing

```
grunt> illustrate path;
```

# illustrate

urls1	A:bytearray	B:bytearray
	url2	url1
	url1	url2
	url5	url1
	url1	url4

urls2	C:bytearray	D:bytearray
	url2	url1
	url1	url2
	url5	url1
	url1	url4

path	urls1::A:bytearray	urls1::B:bytearray	urls2::C:bytearray	urls2::D:bytearray
	url2	url1	url1	url2
	url2	url1	url1	url4
	url5	url1	url1	url2
	url5	url1	url1	url4

# Summary

- Simple database operations translated to Hadoop jobs
- Introduction to Pig

THE END